

# Iterative Best Response of Zero-Sum Racing Game

Clément Suttor

CLEMENT.SUTTOR@EPFL.CH

Yo-Shiun Cheng

YO-SHIUN.CHENG@EPFL.CH

## Abstract

During this project for the course ME-429 MULTIAGENT DECISION-MAKING AND CONTROL, instructed by Dr. Maryam Kamgarpour, we develop a control system in which two autonomous racing cars compete head-to-head. Each car aims to outpace the other by casting their interaction as a zero-sum, multi-stage, dynamic feedback game.

**Keywords:** Iterative best response, zero-sum game, and feedback game.

## 1. Introduction

To make this idea concrete, we simplify each vehicle’s motion using the *bicycle model*, where a single front wheel and a single rear wheel capture acceleration and steering dynamics. At every instant, each car selects an acceleration and a steering angle; these choices affect its own trajectory and, through the shared track and collision constraints, influence the opponent’s possible moves. By planning over a short horizon, solving a quadratic program at each step, and then applying only the first control action before re-planning, both cars continuously adapt to each other—just like real racers glance at each other before deciding whether to overtake or defend.

This formulation is motivated by the need for safe and robust real-time decision making in competitive driving: each autonomous car must anticipate its rival’s tactics, respect track and collision constraints, and still pursue its own racing goals. Our main objective is to merge this optimization framework (quadratic programming) with simple game-theoretic ideas so that, at each time step, the resulting pair of control sequences forms a *saddle point*: neither car can unilaterally improve its outcome.

Game-theoretic planning for competitive autonomy has received growing attention in recent years. Wang *et al.* propose a receding-horizon game-theoretic planner for two-car racing scenarios that represents trajectories as piecewise polynomials, incorporates bicycle kinematics, and adds a sensitivity term to account for collision avoidance and opponent yielding; their Sensitivity-Enhanced IBR algorithm was validated in both numerical simulations and full-scale experiments. More recently, Zhang *et al.* analyze zero-sum linear quadratic (LQ) games as nonconvex–nonconcave saddle-point problems and prove that simple policy-optimization schemes (nested gradient, natural policy gradient, etc.) converge globally to Nash equilibria with sublinear and locally linear rates. In contrast, our work bridges these lines by (i) formulating the two-car racing duel as a zero-sum, multi-stage dynamic feedback game in Frenet coordinates with realistic non-holonomic bicycle dynamics, (ii) solving it online via iterative best–response quadratic programs that enforce track-boundary and collision constraints directly, and (iii) demonstrating convergence properties in simulation.

## 2. Problem Setup

### 2.1. Notation

- $x = [s, d, e_\psi, v]^\top$ : state (longitudinal progress, lateral deviation, heading error, speed).
- $u = [a, \delta]^\top$ : control input (acceleration, steering angle).
- $N$ : prediction horizon (number of steps).
- $\Delta t$ : time step.
- $\alpha_{goal}, \alpha_{opt}, \alpha_{next}, \alpha_u$ : cost weights.
- $s_{goal}, d_{opt_i}, s_{next}$ : longitudinal progress of the target, offset from the optimal trajectory for player  $i$ , longitudinal progress of the next way-point on the optimal trajectory.

### 2.2. Dynamics and Constraints

#### 2.2.1. VEHICLE DYNAMICS

We model each vehicle with a simple bicycle kinematic model, tracking its position, heading, speed, and steering. These continuous- and discrete-time formulations, the Frenet-coordinate transformation, and their linearization are provided in full in the Appendix A.”

#### 2.2.2. CONSTRAINTS

These linear Frenet-frame dynamics are then enforced as constraints in our quadratic programs. At each step  $k$  the Frenet-frame state evolves under the (linearized) discrete dynamics

$$x_{k+1} = A_k x_k + B_k u_{i,k},$$

subject to:

#### State and Input Constraints:

$$\begin{aligned} 0 &\leq v_k \leq v_{\max}, \\ -a_{\max} &\leq a_{i,k} \leq a_{\max}, \\ -\delta_{\max} &\leq \delta_{i,k} \leq \delta_{\max}, \end{aligned}$$

#### Track Boundary Constraint:

$$-\frac{\text{track\_width}}{2} \leq d_k \leq \frac{\text{track\_width}}{2},$$

#### Collision Avoidance Constraint:

Elliptical safety zone

### 2.3. Players, Strategy Sets and Cost Functions

We have two players,  $i \in \{1, 2\}$ :

$$U_i = \{u_{i,0}, u_{i,1}, \dots, u_{i,N-1}\}, \quad u_{i,k} \in \mathcal{U} = [0, a_{\max}] \times [-\delta_{\max}, \delta_{\max}].$$

We split the per-stage cost into four terms:

$$J_{\text{goal}}(s_1, s_2) = \alpha_{\text{goal}} (s_{\text{goal}} - s_1)^2 - \alpha_{\text{goal}} (s_{\text{goal}} - s_2)^2, \quad (1)$$

$$J_{\text{opt}}(d_1, d_2) = \alpha_{\text{opt}} (d_{\text{opt}_1} - d_1)^2 - \alpha_{\text{opt}} (d_{\text{opt}_2} - d_2)^2, \quad (2)$$

$$J_{\text{next}}(s_1, s_2) = \alpha_{\text{next}} (s_{\text{next}} - s_1)^2 - \alpha_{\text{next}} (s_{\text{next}} - s_2)^2, \quad (3)$$

$$J_u(u_1, u_2) = \alpha_u u_1^2 + \alpha_u u_2^2, \quad (4)$$

where  $J_{\text{goal}}$  is the cost w.r.t. the goal,  $J_{\text{opt}}$  is the cost w.r.t. the optimal trajectory,  $J_{\text{next}}$  is the cost w.r.t. the next way point on the optimal trajectory, and  $J_u$  is the cost on control effort.

The stage cost for Player 1 is then

$$J = J_{\text{goal}}(s_1, s_2) + J_{\text{opt}}(d_1, d_2) + J_{\text{next}}(s_1, s_2) + J_u(u_1, u_2)$$

and the total cost is

$$\mathcal{J}_1 = \sum_{k=0}^{N-1} J_k + J_{\text{terminal}}.$$

Player 1 seeks to *minimize* a total cost  $\mathcal{J}_1$ ; Player 2 seeks to *maximize* the same (so  $\mathcal{J}_2 = -\mathcal{J}_1$ ). This makes the game zero-sum.

## 2.4. Game Classification

- *Zero-sum*:  $\mathcal{J}_2 = -\mathcal{J}_1$ .
- *Multi-stage*, because decisions are made for  $k = 0, \dots, N - 1$ .
- *Dynamic*, since state updates couple decisions over time.
- *Feedback*, as both players re-solve at each  $k$  based on the current  $x_k$ .

## 3. Analysis

### 3.1. Iterative Best-Response Algorithm

To compute the saddle-point controls at each time step, we employ an *iterative best-response* scheme. Starting from an initial guess for the opponent's trajectory, each player alternately solves a quadratic program (QP) to optimize its own cost while holding the other fixed. We repeat until neither player can significantly improve. (Details shown in Appendix B.)

### 3.2. Experimental Setup

All code is written in Python and organized into the following modules:

- `car.py` / `car_dynamics.py`: defines the `Car` class and the RK4-discretized bicycle model dynamics. (Appendix C and D)
- `frenet.py`: routines for computing Frenet coordinates  $(s, d)$  relative to a spline-interpolated centerline. (Appendix F)

- `optimal_trajectory.py`: generates a smooth, time-optimal lateral offset  $d_{\text{opt}}(s)$  via `scipy.optimize.minimize` (Appendix G)
- `racetrack.py`: builds an “U-shaped” track with the `Track` class and solves for waypoints and headings. (Appendix E)
- `ZS.Controller.py`: implements the zero-sum best-response QPs for both cars using `cvxpy` and OSQP. (Appendix H)

### 3.2.1. SIMULATION WORKFLOW

To put all the things together, this is how we do it:

**Track Generation** We call

```
track = generate_track (Length=100, Width=6)
```

to create an U-shaped track. We extract `centerline = [xm, ym]` and `headings = θm` from the `Track` object, then build the arc-length vector `centerline_frenet` by cumulative sum of Euclidean distances.

**Optimal Trajectories** For each car  $i \in \{1, 2\}$  we compute

```
dopti = generate_optimal_traj(centerline_frenet, track.width, alat_maxi)
```

which yields a smooth lateral offset that approximately minimizes lap time.

**Car Initialization** Two `Car` objects are created with different maximum acceleration and maximum velocity to break symmetry:

```
car1 = Car([x0, y0 + 2, θ0, 5.0], vmax = 75, amax = 10, alat_max = g),
car2 = Car([x0, y0 - 2, θ0, 5.0], vmax = 65, amax = 12, alat_max = 2g),
```

where  $(x_0, y_0, \theta_0)$  is the start of the centerline.

**Closed-Loop Simulation** We set

$$\Delta t = 0.05 \text{ s}, \quad N = 20, \quad T = \frac{10 \text{ s}}{\Delta t} = 200 \text{ steps},$$

At each step  $t$ :

1. Compute Frenet positions  $s_i, d_i$  for both cars.
2. Determine next reference  $s_{\text{next}_i}$  by advancing one waypoint.
3. Call

```
u1, u2 = zero_sum_best_response(..., N, Δt, ...),
```

which returns the first control inputs for each car.

4. Apply `Car.set_control` and `Car.update`.

**Plotting** After simulation, we plot the track boundaries and the two car histories using `matplotlib.pyplot`.

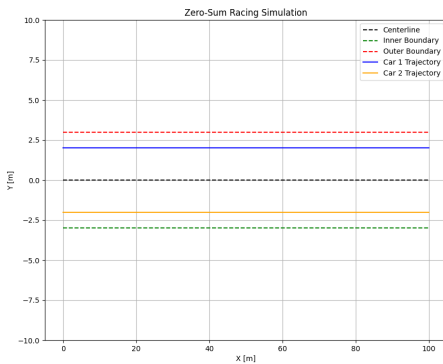
#### 4. Simulation Results

Despite extensive attempts, we were only able to obtain viable results for a **straight-line track**. We experimented with modifying the framework from Cartesian to Frenet coordinates in the Quadratic Program (QP), adjusting the cost function, and relaxing constraints. Although relaxing the constraints enabled the solver to compute a solution, the resulting vehicle behavior was incorrect. Nevertheless, the straight-line case allows us to extract meaningful insights about our implementation.


Figure 1(a) shows the trajectories of the two cars. As expected, both vehicles accelerate forward at maximum speed toward the finish line which is an optimal behavior in this simplified setting. This confirms that the saddle-point controller successfully computed optimal inputs under these conditions.

More interestingly, we can use this scenario to analyze the convergence of our algorithm toward a *saddle-point equilibrium* in the zero-sum dynamic game. According to [2], unconstrained policy optimization converges to equilibrium for a certain class of linear-quadratic dynamic games. We wanted to see if these results applied to our best-response algorithm despite having constraints in the QPs. Figure 1(b) shows a boxplot of the the cost difference  $|J_1 - J_2|$  distribution over timesteps of the simulation. The mean and median cost differences are 296.9 and 269.4, respectively. Given that the average total cost per player is on the order of  $\sim 3000$ , this implies a relative cost gap  $|\bar{V} - \underline{V}|$  of only  $\sim 10\%$ .

Combined with the fact that the computed control inputs align with the expected optimal behavior, these results suggest that our best-response algorithm approximates well the saddle-point equilibrium at each time step in this scenario.



(a) Car Trajectories on the Track



(b) Distribution of  $|J_1 - J_2|$  across timesteps

Figure 1: Simulation Results for the Straight Line Track

## 5. Conclusion

In this project, we wanted to address the problem of real-time decision-making for autonomous vehicles engaged in competitive racing, formulated as a zero-sum, multi-stage dynamic feedback game. Our goal was to merge game-theoretic planning with optimization techniques to enable two vehicles to anticipate each other’s strategies and make rational, competitive driving decisions.

We successfully developed a simulation framework using the iterative best-response approach, applied to cars modeled by bicycle dynamics in Frenet coordinates. The optimal control strategies were derived from quadratic programs that enforced track boundaries and collision avoidance constraints on top of the usual state and input constraints. While the controller struggled to converge on more complex tracks, the straight-line case provided critical validation: both vehicles exhibited optimal behavior and more importantly our algorithm demonstrated convergence toward saddle-point equilibrium solutions.

These results demonstrate that iterative best-response (IBR) methods can effectively approximate saddle-point equilibria in constrained zero-sum dynamic games. Nevertheless, several important questions remain open. Notably, the observed instability of the framework on complex, curved tracks suggests a need for more robust formulations. Improvements may involve refining the constraint handling, enhancing the linearization of vehicle dynamics, or adopting more sophisticated numerical solvers to ensure convergence under tighter feasibility conditions.

Moreover, the modular nature of the proposed framework allows for the investigation and comparison of alternative solutions within the same zero-sum game-theoretic setting. For instance, a natural extension would be to implement a Stackelberg (leader–follower) formulation, where one agent commits to a strategy that the other responds to optimally.

Investigating these alternative strategies may yield valuable insights into robustness, computational efficiency, and applicability across more diverse and realistic racing scenarios.

## 6. Partnership work and resources

We contributed equally, moving from a grid-based zero-sum racing setup to a continuous-track formulation to tackle collision handling. We briefly split—Clément refining the continuous controller and Yo-Shiun building a first-price auction backup project—then reunited under TA guidance for debugging.

We’ve leveraged AI throughout our workflow to accelerate debugging and enhance writing: by feeding error messages and solver failure codes into a large language model, we iteratively refined our vehicle dynamics and QP formulations, rapidly identifying linearization mistakes and solver parameter tweaks. Additionally, AI-powered translation tools seamlessly converted Chinese drafts into polished English, ensuring clear documentation among all the project.

## **Acknowledgments**

We'd like to thank our TAs for their invaluable insights and guidance throughout this project, which helped us refine our approach and stay on track.

## **References**

- [1] Mingyu Wang et al. "Game Theoretic Planning for Self-Driving Cars in Competitive Scenarios".
- [2] Kaiqing Zhang, Zhuoran Yang, and Tamer Basar. "Policy optimization provably converges to Nash equilibria in zero-sum linear quadratic games".

## Appendix A. Vehicle Dynamics

To describe each car's motion, we use the *bicycle model*, which reduces the four wheels to a single front wheel and a single rear wheel. Let the state vector at time  $t$  be

$$x(t) = \begin{bmatrix} x(t) \\ y(t) \\ \theta(t) \\ v(t) \end{bmatrix},$$

where

- $x, y$  are the car's position coordinates in the plane,
- $\theta$  is the heading angle,
- $v$  is the forward speed.

The control inputs are

$$u(t) = \begin{bmatrix} a(t) \\ \delta(t) \end{bmatrix},$$

where  $a$  is the longitudinal acceleration and  $\delta$  is the steering angle of the front wheel.

### Continuous-Time Dynamics

The continuous-time equations of motion are

$$\dot{x} = v \cos \theta, \tag{5}$$

$$\dot{y} = v \sin \theta, \tag{6}$$

$$\dot{\theta} = \frac{v}{L} \tan \delta, \tag{7}$$

$$\dot{v} = a, \tag{8}$$

where  $L$  is the wheelbase (distance between front and rear wheel centers).

### Frenet Coordinate Transformation

To simplify planning along a curved track, we express the vehicle state in *Frenet coordinates*  $(s, d)$ , where

- $s$  is the longitudinal distance along a reference path  $r(s) = (x_r(s), y_r(s))$ ,
- $d$  is the signed lateral offset from that path,
- $\theta_r(s)$  is the tangent angle of the path at  $s$ ,
- $\kappa_r(s)$  is the curvature of the path at  $s$ .



Define the heading error

$$e_\psi = \theta - \theta_r(s).$$

Then the continuous-time dynamics in Frenet frame become:

$$\dot{s} = \frac{v \cos(e_\psi)}{1 - \kappa_r(s) d}, \quad (9)$$

$$\dot{d} = v \sin(e_\psi), \quad (10)$$

$$\dot{e}_\psi = \frac{v}{L} \tan \delta - \frac{\kappa_r(s) v \cos(e_\psi)}{1 - \kappa_r(s) d}, \quad (11)$$

$$\dot{v} = a. \quad (12)$$

Here, the denominator  $1 - \kappa_r(s) d$  accounts for the fact that moving off the centerline effectively shortens (or lengthens) the path curvature. These equations define

$$\dot{x} = f(x, u),$$

with

$$x = \begin{bmatrix} s \\ d \\ e_\psi \\ v \end{bmatrix}, \quad u = \begin{bmatrix} a \\ \delta \end{bmatrix}.$$

### Discretization and Linearization in Frenet Frame

We discretize  $\dot{x} = f(x, u)$  with RK4 to obtain

$$x_{k+1} = g(x_k, u_k).$$

Then around the current operating point  $(x_k, u_k)$  we form the linear approximation

$$x_{k+1} \approx A_k x_k + B_k u_k,$$

where

$$A_k = \left. \frac{\partial g}{\partial x} \right|_{(x_k, u_k)}, \quad B_k = \left. \frac{\partial g}{\partial u} \right|_{(x_k, u_k)}.$$

## Appendix B. Algorithm Description

---

### Algorithm 1 Iterative Best-Response at time step $k$

---

**Require:** current state  $x_{f,k}$

1: Initialize

$$U_2^{(0)} \leftarrow \text{zero}, \quad U_1^{(0)} \leftarrow \text{zero}.$$

2: **for**  $i = 1, \dots, i_{\max}$  **do**

3:   **Player 1 (minimizer) update:**

$$U_1^{(i)} = \arg \min_{U_1} J_1(U_1, U_2^{(i-1)}) \quad \text{s.t. dynamics \& constraints}$$

4:   Warm-start QP solver with  $U_1^{(i-1)}$

5:   **Player 2 (maximizer) update:**

$$U_2^{(i)} = \arg \max_{U_2} J_1(U_1^{(i)}, U_2) = \arg \min_{U_2} -J_1(U_1^{(i)}, U_2) \quad \text{s.t. dynamics \& constraints}$$

6:   Warm-start QP solver with  $U_2^{(i-1)}$

7:   **if**  $\|U_1^{(i)} - U_1^{(i-1)}\|_{\infty} < \epsilon$  **and**  $\|U_2^{(i)} - U_2^{(i-1)}\|_{\infty} < \epsilon$  **then**

8:     **break**

9:   **end if**

10: **end for**

11: **Return**  $U_1^* = U_1^{(i)}, U_2^* = U_2^{(i)}$

---

### Details and Parameters

- $\epsilon$  is the convergence tolerance (e.g.  $10^{-6}$ ).
- $i_{\max}$  is a safeguard on iterations (e.g. 5).
- *Warm-starting* the QP solvers with the previous iterate greatly reduces solve time.
- Once  $(U_1^*, U_2^*)$  is found, only the first control inputs  $(u_{1,k}^*, u_{2,k}^*)$  are applied; then the process repeats at the next time step  $k + 1$ .

This procedure ensures that at each discretized time  $k$ , the pair  $(U_1^*, U_2^*)$  approximates the saddle-point of the finite-horizon game, yielding a feedback policy that continuously reacts to the opponent's moves.

## Appendix C. Car Class Code

car.py

```
1 import numpy as np
2 from car_dynamics import bicycle_dynamics_discrete
```

```

3 from frenet import get_frenet_coords
4
5 class Car:
6     def __init__(self, initial_state, v_max, a_max, a_lat_max,
7         max_steering_angle=np.pi/2, dt=0.1, L=2.5):
8         self.state = np.array(initial_state, dtype=float)
9         self.control_inputs = np.array([0.0, 0.0])
10        self.dt = dt
11        self.L = L
12        self.history = [self.state.copy()]
13        self.v_max = v_max
14        self.a_max = a_max
15        self.a_lat_max = a_lat_max
16        self.max_steering_angle = max_steering_angle
17
18    def set_control(self, a, delta):
19        """Set current control inputs."""
20        self.control_inputs = np.array([a, delta])
21
22    def update(self):
23        """Advance the state using bicycle dynamics."""
24        self.state = bicycle_dynamics_discrete(self.state, self.
25            control_inputs, self.dt, self.L)
26        self.history.append(self.state.copy())
27
28    def get_state(self):
29        """Return the current state."""
30        return self.state
31
32    def get_control_inputs(self):
33        """Return the current control inputs."""
34        return self.control_inputs
35
36    def get_frenet_coords(self, centerline, headings):
37        """Return current (s, d) in Frenet frame."""
38        x, y = self.state[0], self.state[1]
39        return get_frenet_coords(x, y, centerline, headings)
40
41    def compute_delta_psi(self, centerline, headings):
42        """Compute delta_psi for frenet dynamics"""
43        x, y, theta = self.state[0], self.state[1], self.state[2]
44        diffs = centerline - np.array([x, y])
45        idx = np.argmin(np.sum(diffs**2, axis=1))
46        psi_ref = headings[idx]
47        delta_psi = theta - psi_ref
48        delta_psi = (delta_psi + np.pi) % (2 * np.pi) - np.pi
49        return delta_psi

```

## Appendix D. Car Dynamics Code

car\_dynamics.py

```

1 import numpy as np
2
3 def bicycle_dynamics_continuous(state, control, L=2.5):
4     """Continuous-time bicycle model dynamics."""
5     x, y, theta, v = state
6     a, delta = control
7
8     x_dot = v * np.cos(theta)
9     y_dot = v * np.sin(theta)
10    theta_dot = v / L * np.tan(delta)
11    v_dot = a
12
13    return np.array([x_dot, y_dot, theta_dot, v_dot])
14
15
16 def frenet_bicycle_dynamics(x, u, kappa, L=2.5):
17     """Continuous-time bicycle model dynamics in Frenet frame."""
18     s, d, v, delta_psi = x
19     a, delta = u
20
21     denom = max(1.0 - kappa * d, 1e-5)
22
23     s_dot = v * np.cos(delta_psi) / denom
24     d_dot = v * np.sin(delta_psi)
25     v_dot = a
26     delta_psi_dot = v / L * np.tan(delta) - kappa * v * np.cos(delta_psi)
27         / denom
28
29     return np.array([s_dot, d_dot, v_dot, delta_psi_dot])
30
31
32 def rk4_integration(state, control, dt, dynamics):
33     """RK4 integration step."""
34     k1 = dynamics(state, control)
35     k2 = dynamics(state + 0.5 * dt * k1, control)
36     k3 = dynamics(state + 0.5 * dt * k2, control)
37     k4 = dynamics(state + dt * k3, control)
38
39     return state + (dt / 6.0) * (k1 + 2*k2 + 2*k3 + k4)
40
41
42 def bicycle_dynamics_discrete(state, control, dt, L=2.5):
43     """Bicycle model dynamics discretized using RK4"""
44     x_plus = rk4_integration(state, control, dt, lambda s, u:
45         bicycle_dynamics_continuous(s, u, L))
46     return x_plus

```

```

47 def frenet_dynamics_discrete(x, u, kappa_func, dt, L=2.5):
48     """Frenet bicycle dynamics discretized using RK4"""
49     def dynamics(x_local, u_local):
50         s = x_local[0]
51         kappa = kappa_func(s)
52         return frenet_bicycle_dynamics(x_local, u_local, kappa, L)
53     x_plus = rk4_integration(x, u, dt, dynamics)
54     return x_plus

```

## Appendix E. Track Generation Code

racetrack.py

```

1  from trackgen import Track
2  from math import pi
3  import numpy as np
4
5  def generate_track(Length = 1000, Width = 6):
6      """ Generates the track following the example template provided on
7          the github of the trackgen package : https://github.com/mopg/trackgen"""
8      # ## Oval
9      # # Where are the corners?
10     # # crns = np.array( [False,True,False,True], dtype=bool )
11     # crns = np.array( [False,True,False], dtype=bool )
12
13     # # Change in angle (needs to be zero for straight)
14     # # delTh = np.array( [0,pi,0,pi], dtype=float )
15     # delTh = np.array( [0,pi/2,0], dtype=float )
16
17     # # length parameter initial guess (radius for corner, length for
18     # # straight)
19     # # lpar = np.array( [Length/2,Length/5,Length/2,Length/5], dtype=
20     # # float )
21     # lpar = np.array( [Length/2,Length/5,Length/2], dtype=float )
22
23     # # Solve
24     # track = Track( length = Length, width = Width, left = True, crns =
25     # # crns )
26     # sol = track.solve( lpar, delTh, case = 2 )
27
28     ## Straight line
29     track = Track(
30         width = Width,
31         crns = np.array([False]),
32         lpar = np.array([Length]),
33         delTh = np.array([0.0]),
34     )
35     track.compTrackXY()
36

```

```

33     xe, ye, thcum = track.endpoint()
34
35     print( "End point    = (%4.3f, %4.3f)" % (xe, ye) )
36     print( "Final angle =  %4.3f" % (thcum) )
37
38     return track

```

## Appendix F. Usefull Fuctions in the Frenet Framework

frenet.py

```

1  import numpy as np
2
3  def get_frenet_coords(x, y, centerline, headings):
4      '''Computes the Frenet coordinates (s,d) of a point in cartesian
5         coordinates (x,y) w.r.t to the centerline'''
6      # Compute closest point on the centerline
7      dx = centerline[:,0] - x
8      dy = centerline[:,1] - y
9      distances = np.hypot(dx,dy)
10     idx = np.argmin(distances)
11
12     # Compute arc length s
13     s = np.sum(np.hypot(np.diff(centerline[:idx+1, 0]), np.diff(
14         centerline[:idx+1, 1])))
15
16     # Compute deviation d
17     path_heading = headings[idx]
18     normal = np.array([-np.sin(path_heading), np.cos(path_heading)])
19     rel_pos = np.array([x, y]) - centerline[idx]
20     d = np.dot(rel_pos, normal)
21
22     return s, d
23
24 def compute_curvature(centerline):
25     '''Computes the curvature Kappa along the centerline'''
26     x = centerline[:, 0]
27     y = centerline[:, 1]
28
29     dx = np.gradient(x)
30     dy = np.gradient(y)
31     ddx = np.gradient(dx)
32     ddy = np.gradient(dy)
33
34     num = dx * ddy - dy * ddx
35     denom_raw = (dx**2 + dy**2) ** 1.5
36     # Avoid division by 0
37     denom = np.where(np.abs(denom_raw) < 1e-6, 1e-6, denom_raw)
38     kappa = np.divide(num, denom)
39     return kappa

```

## Appendix G. Optimal Trajectory Computation Code

optimal\_trajectory.py

```

1  import numpy as np
2  from scipy.optimize import minimize
3
4  def lap_time_cost_with_smoothness(d, s, a_lat_max=9.0, w_smooth=0.5):
5      """Computes the cost that is used in the optimization problem to
6          compute the optimal lateral offsets from the centerline."""
7      # Estimation of the lap time
8      ds = np.gradient(s)
9      d2_ds2 = np.gradient(np.gradient(d, s), s)
10     curvature = np.abs(d2_ds2)
11     curvature = np.clip(curvature, 1e-4, None)
12
13     v_max = np.sqrt(a_lat_max / curvature)
14     segment_times = ds / v_max
15     lap_time = np.sum(segment_times)
16
17     # Penalize fast changes in lateral position to have a smooth optimal
18     # trajectory
19     smoothness = np.sum(np.diff(d, 2)**2)
20
21     total_cost = lap_time + w_smooth * smoothness
22
23     return total_cost
24
25 def generate_optimal_traj(centerline_frenet, track_width, a_lat_max=9.0)
26 :
27     """Computes the optimal lateral offsets from the centerline to
28     minimize lap time."""
29     N = len(centerline_frenet)
30     max_dev = track_width / 2 * 0.95
31     np.random.seed(42)
32     initial_d = np.random.randn(N) * max_dev
33     bounds = [(-max_dev, max_dev)] * N # Ensures that the optimal
34     lateral offsets remain inside of the track
35
36     res = minimize(
37         lap_time_cost_with_smoothness,
38         initial_d,
39         args=(centerline_frenet, a_lat_max),
40         bounds=bounds,
41         method='L-BFGS-B'
42     )
43
44     return res.x

```

## Appendix H. Best Iterative Response Controller Code

ZS\_controller.py

```

1  import numpy as np
2  import cvxpy as cp
3  from car import Car
4  from car_dynamics import bicycle_dynamics_discrete
5  from car_dynamics import frenet_dynamics_discrete
6  from frenet import get_frenet_coords
7  from scipy.signal import cont2discrete
8  from scipy.interpolate import interp1d
9
10 def linearize_bicycle_dynamics(x_ref, u_ref, L=2.5):
11     '''Computes the A and B matrices corresponding to the linearized
12         bicycle dynamics around (x_ref, u_ref)'''
13     _, _, theta, v = x_ref
14     a, delta = u_ref
15
16     A = np.zeros((4, 4))
17     B = np.zeros((4, 2))
18
19     # Compute jacobian w.r.t to state to get A
20     A[0, 2] = -v * np.sin(theta)
21     A[0, 3] = np.cos(theta)
22     A[1, 2] = v * np.cos(theta)
23     A[1, 3] = np.sin(theta)
24     A[2, 3] = 1.0 / L * np.tan(delta)
25
26     # Compute jacobian w.r.t to input to get B
27     B[2, 1] = v / L / (np.cos(delta)**2)
28     B[3, 0] = 1.0
29
30     return A, B
31
32 def linearize_frenet_dynamics(x_ref, u_ref, kappa, L=2.5):
33     '''Computes the A and B matrices corresponding to the linearized
34         bicycle dynamics in frenet coordinates around (x_ref, u_ref,
35         kappa_ref)'''
36     s, d, v, delta_psi = x_ref
37     a, delta = u_ref
38
39     # Compute constants used in the jacobian
40     denom = 1.0 - kappa * d
41     denom = np.clip(denom, 1e-6, None) # prevent division by 0
42     cos_dpsi = np.cos(delta_psi)
43     sin_dpsi = np.sin(delta_psi)
44     cos_dpsi = np.clip(cos_dpsi, -1.0, 1.0)
45
46     A = np.zeros((4, 4))
47     B = np.zeros((4, 2))

```



## ZERO-SUM RACING GAME

```

46     # Compute jacobian w.r.t to state to get A
47     A[0, 1] = v * kappa * cos_dpsi / (denom ** 2)
48     A[0, 2] = cos_dpsi / denom
49     A[0, 3] = -v * sin_dpsi / denom
50     A[1, 2] = sin_dpsi
51     A[1, 3] = v * cos_dpsi
52     A[3, 1] = -v * kappa**2 * cos_dpsi / (denom ** 2)
53     A[3, 2] = delta / L - kappa * cos_dpsi / denom
54     A[3, 3] = v * kappa * sin_dpsi / denom
55
56     # Compute jacobian w.r.t to input to get B
57     B[2, 0] = 1.0
58     B[3, 1] = v / (L * (np.cos(delta) ** 2))
59
60     return A, B
61
62 def discretize_linear_dynamics(A, B, dt):
63     '''Discretize matrices A and B given dt using Zero-Order-Hold'''
64     C = np.eye(A.shape[0])
65     D = np.zeros(B.shape)
66     system = (A, B, C, D)
67     A_d, B_d, _, _, _ = cont2discrete(system, dt, method='zoh')
68     return A_d, B_d
69
70
71 ## Solve the problem in (x,y) Coordinates
72 def solve_zero_sum_qp_P1(
73     car1, car2,
74     d_opt1, d_opt2, centerline_frenet,
75     centerline, headings, track_width, s_goal, s_next_1, s_next_2,
76     P2_control_input,
77     N=10, dt=0.1, d_g=1.0, d_u=1.0, d_t=0.2, d_next=2.0, d_c=100.0
78 ):
79     n, m = 4, 2
80     X1 = cp.Variable((n, N+1))
81     U1 = cp.Variable((m, N))
82     U2 = P2_control_input
83     slack = cp.Variable(N)
84     x1_0 = car1.get_state()
85     x2_0 = car2.get_state()
86     s1_0, d1_0 = car1.get_frenet_coords(centerline, headings)
87
88     constraints = [
89         X1[:, 0] == x1_0,
90     ]
91
92     cost = 0.0
93
94     # Assume horizon is short enough to linearize around x_0
95     A1, B1 = linearize_bicycle_dynamics(x1_0, [0.0, 0.0])
96     A1_d, B1_d = discretize_linear_dynamics(A1, B1, dt)

```

```

97
98 # Create interpolators locally
99 d_opt_interp1 = interp1d(centerline_frenet, d_opt1, kind='linear',
100    fill_value='extrapolate')
101
102 d_opt_interp2 = interp1d(centerline_frenet, d_opt2, kind='linear',
103    fill_value='extrapolate')
104
105 # Compute desired lateral offset
106 d_goal1 = d_opt_interp1(s1_0)
107
108 # Compute parameters for frenet coordinates linearization of car 1
109 idx1 = np.argmin(np.abs(centerline_frenet - s1_0))
110 theta_c1 = headings[idx1]
111 t_hat1 = np.array([np.cos(theta_c1), np.sin(theta_c1)])
112 n_hat1 = np.array([-np.sin(theta_c1), np.cos(theta_c1)])
113 p1_0 = x1_0[0:2]
114
115 # Rotation matrix for whole horizon (assume horizon short enough for
116    the cars to keep the same orientation during whole length)
117 theta_rel = 0.5 * (x1_0[2] + x2_0[2]) # average heading of the cars
118 R = np.array([[np.cos(theta_rel), np.sin(theta_rel)],
119    [-np.sin(theta_rel), np.cos(theta_rel)]])
120
121 # Initialize state variables for car 2
122 x2_k = x2_0
123
124 for k in range(N):
125
126     # Constraints on dynamics
127     constraints += [X1[:, k+1] == A1_d @ X1[:, k] + B1_d @ U1[:, k]]
128
129     # Box constraints on speed, acceleration, and steering
130     constraints += [X1[3, k] <= car1.v_max, X1[3, k] >= 0.0]
131     constraints += [U1[0, k] <= car1.a_max, U1[0, k] >= -car1.a_max]
132     constraints += [U1[1, k] <= car1.max_steering_angle, U1[1, k] >=
133        -car1.max_steering_angle]
134
135     # Collision constraints using an ellipse rotated in the
136        direction of the cars (both cars cannot be in the ellipse at
137        the same time)
138     delta = X1[0:2, k] - x2_k[0:2]
139     rel_rot = R @ delta
140     collision_expr = cp.quad_form(rel_rot, np.diag([1/2.0**2,
141        1/1.0**2]))
142     constraints += [collision_expr + slack[k] >= 1.0]
143     cost += d_c * slack[k]
144
145     # Slack variables always positive
146     constraints += [slack[k] >= 0]
147
148     # Linear approximation of frenet coordinates for X1[0:2,k]

```

## ZERO-SUM RACING GAME

```

141     s1 = s1_0 + t_hat1 @ (X1[0:2, k] - p1_0)
142     d1 = d1_0 + n_hat1 @ (X1[0:2, k] - p1_0)
143
144     # Constraints on lateral deviation to prevent the car from going
145     # outside the track
146     constraints += [d1 <= track_width/2*0.95, d1 >= -track_width
147                     /2*0.95]
148
149     # Use real frenet coordiantes for car 2
150     s2, d2 = get_frenet_coords(x2_k[0], x2_k[1], centerline,
151                               headings)
152
153     # Update optimal trajectory goal
154     d_goal2 = d_opt_interp2(s2)
155
156     # Goal cost
157     cost += d_g * cp.square(s_goal - s1) - d_g * cp.square(s_goal -
158                     s2)
159
160     # Optimal trajectory cost
161     cost += d_t * cp.square(d_goal1 - d1) - d_t * cp.square(d_goal2
162                     - d2)
163
164     # Input cost on acceleration
165     cost += d_u * cp.square(U1[0, k]) - d_u * cp.square(U2[0, k])
166
167     # Next waypoint cost
168     cost += d_next * cp.square(s_next_1 - s1) - d_next * cp.square(
169                     s_next_2 - s2)
170
171     # Update state of car 2
172     x2_k = bicycle_dynamics_discrete(x2_k, U2[:,k], dt)
173
174     # Linear approximation of frenet coordinates for X[0:2,N]
175     s1_N = s1_0 + t_hat1 @ (X1[0:2, N] - p1_0)
176     d1_N = d1_0 + n_hat1 @ (X1[0:2, N] - p1_0)
177
178     # Real frenet coordinates for X2_N
179     s2_N, d2_N = get_frenet_coords(x2_k[0], x2_k[1], centerline,
180                               headings)
181
182     # Goal cost
183     cost += d_g * cp.square(s_goal - s1_N) - d_g * cp.square(s_goal -
184                     s2_N)
185
186     # Optimal trajectory cost
187     cost += d_t * cp.square(d_goal1 - d1_N) - d_t * cp.square(d_goal2 -
188                     d2_N)
189
190     # Next waypoint cost

```

```

182     cost += d_next * cp.square(s_next_1 - s1_N) - d_next * cp.square(
183         s_next_2 - s2_N)
184
185     # Problem solving
186     prob = cp.Problem(cp.Minimize(cost), constraints)
187     prob.solve(solver=cp.OSQP)
188
189     # Check if solver succeeded
190     if prob.status not in [cp.OPTIMAL, cp.OPTIMAL_INACCURATE]:
191         print("[Warning] Car 1 QP solver failed:", prob.status)
192         return np.zeros((2, N))
193
194     return U1.value, prob.value
195
196 def solve_zero_sum_qp_P2(
197     car1, car2,
198     d_opt1, d_opt2, centerline_frenet,
199     centerline, headings, track_width, s_goal, s_next_1, s_next_2,
200     P1_control_input,
201     N=10, dt=0.1, d_g=1.0, d_u=1.0, d_t=0.2, d_next=2.0, d_c=100.0
202 ):
203     n, m = 4, 2
204     X2 = cp.Variable((n, N+1))
205     U2 = cp.Variable((m, N))
206     U1 = P1_control_input
207     slack = cp.Variable(N)
208     x1_0 = car1.get_state()
209     x2_0 = car2.get_state()
210     s2_0, d2_0 = car2.get_frenet_coords(centerline, headings)
211
212     constraints = [
213         X2[:, 0] == x2_0,
214     ]
215
216     cost = 0.0
217
218     # Assume horizon is short enough to linearize around x_0
219     A2, B2 = linearize_bicycle_dynamics(x2_0, [0.0, 0.0])
220     A2_d, B2_d = discretize_linear_dynamics(A2, B2, dt)
221
222     # Create interpolators locally
223     d_opt_interp1 = interp1d(centerline_frenet, d_opt1, kind='linear',
224         fill_value='extrapolate')
225     d_opt_interp2 = interp1d(centerline_frenet, d_opt2, kind='linear',
226         fill_value='extrapolate')
227
228     # Compute desired lateral offset
229     d_goal2 = d_opt_interp2(s2_0)
230
231     # Compute parameters for frenet coordinates linearization of car 2
232     idx2 = np.argmin(np.abs(centerline_frenet - s2_0))

```

```

230 theta_c2 = headings[idx2]
231 t_hat2 = np.array([np.cos(theta_c2), np.sin(theta_c2)])
232 n_hat2 = np.array([-np.sin(theta_c2), np.cos(theta_c2)])
233 p2_0 = x2_0[0:2]
234
235 # Rotation matrix for whole horizon (assume horizon short enough for
    the cars to keep the same orientation during whole length)
236 theta_rel = 0.5 * (x1_0[2] + x2_0[2]) # average heading of the cars
237 R = np.array([[np.cos(theta_rel), np.sin(theta_rel)],
    [-np.sin(theta_rel), np.cos(theta_rel)]])
238
239 # Initialize state variables for car 1
240 x1_k = x1_0
241
242 for k in range(N):
243
244     # Constraints on dynamics
245     constraints += [X2[:, k+1] == A2_d @ X2[:, k] + B2_d @ U2[:, k]]
246
247     # Box constraints on speed, acceleration, and steering
248     constraints += [X2[3, k] <= car2.v_max, X2[3, k] >= 0.0]
249     constraints += [U2[0, k] <= car2.a_max, U2[0, k] >= -car2.a_max]
250     constraints += [U2[1, k] <= car2.max_steering_angle, U2[1, k] >=
    -car2.max_steering_angle]
251
252     # Collision constraints using an ellipse rotated in the
    direction of the cars (both cars cannot be in the ellipse at
    the same time)
253     delta = x1_k[0:2] - X2[0:2, k]
254     rel_rot = R @ delta
255     collision_expr = cp.quad_form(rel_rot, np.diag([1/2.0**2,
    1/1.0**2]))
256     constraints += [collision_expr + slack[k] >= 1.0]
257     cost += d_c * slack[k]
258
259     # Slack variables always positive
260     constraints += [slack[k] >= 0]
261
262     # Linear approximation of frenet coordinates for X2[0:2,k]
263     s2 = s2_0 + t_hat2 @ (X2[0:2, k] - p2_0)
264     d2 = d2_0 + n_hat2 @ (X2[0:2, k] - p2_0)
265
266     # Constraints on lateral deviation to prevent the car from going
    outside the track
267     constraints += [d2 <= track_width/2*0.95, d2 >= -track_width
    /2*0.95]
268
269     # Use real frenet coordiantes for car 2
270     s1, d1 = get_frenet_coords(x1_k[0], x1_k[1], centerline,
    headings)
271
272
    
```

```

273
274     # Update optimal trajectory goal
275     d_goal1 = d_opt_interp1(s1)
276
277     # Goal cost
278     cost += d_g * cp.square(s_goal - s1) - d_g * cp.square(s_goal -
279         s2)
280
281     # Optimal trajectory cost
282     cost += d_t * cp.square(d_goal1 - d1) - d_t * cp.square(d_goal2
283         - d2)
284
285     # Input cost on acceleration
286     cost += d_u * cp.square(U1[0, k]) - d_u * cp.square(U2[0, k])
287
288     # Next waypoint cost
289     cost += d_next * cp.square(s_next_1 - s1) - d_next * cp.square(
290         s_next_2 - s2)
291
292     # Update state of car 1
293     x1_k = bicycle_dynamics_discrete(x1_k, U1[:,k], dt)
294
295     # Linear approximation of frenet coordinates for X2[0:2,N]
296     s2_N = s2_0 + t_hat2 @ (X2[0:2, N] - p2_0)
297     d2_N = d2_0 + n_hat2 @ (X2[0:2, N] - p2_0)
298
299     # Real frenet coordinates for X2_N
300     s1_N, d1_N = get_frenet_coords(x1_k[0], x1_k[1], centerline,
301         headings)
302
303     # Goal cost
304     cost += d_g * cp.square(s_goal - s1_N) - d_g * cp.square(s_goal -
305         s2_N)
306
307     # Optimal trajectory cost
308     cost += d_t * cp.square(d_goal1 - d1_N) - d_t * cp.square(d_goal2 -
309         d2_N)
310
311     # Next waypoint cost
312     cost += d_next * cp.square(s_next_1 - s1_N) - d_next * cp.square(
313         s_next_2 - s2_N)
314
315     # Solve the problem
316     prob = cp.Problem(cp.Minimize(-cost), constraints)
317     prob.solve(solver=cp.OSQP)
318
319     # Check if solver succeeded
320     if prob.status not in [cp.OPTIMAL, cp.OPTIMAL_INACCURATE]:
321         print("[Warning] Car 2 QP solver failed:", prob.status)
322         return np.zeros((2, N))

```

```

317     return U2.value, prob.value
318
319 def zero_sum_best_response(
320     car1, car2,
321     d_opt1, d_opt2, centerline_frenet,
322     centerline, headings, track_width, s_goal, s_next_1, s_next_2,
323     N=10, dt=0.1, d_g=1.0, d_u=1.0, d_t=20.0, d_next=20.0, d_c=100.0,
324     max_iters=5, epsilon=1e-6
325 ):
326     np.random.seed(42)
327     m = 2
328     U1 = np.zeros((m, N))
329     U2 = np.zeros((m, N))
330
331     for i in range(max_iters):
332         print(f" Iter {i+1}/{max_iters}")
333
334         # Minimize cost for Car 1 given U2
335         U1_new, cost_J1 = solve_zero_sum_qp_P1(
336             car1, car2,
337             d_opt1, d_opt2, centerline_frenet,
338             centerline, headings, track_width, s_goal, s_next_1,
339             s_next_2,
340             U2, N, dt, d_g, d_u, d_t, d_next, d_c
341         )
342
343         # Maximize cost for Car 2 given U1
344         U2_new, cost_J2 = solve_zero_sum_qp_P2(
345             car1, car2,
346             d_opt1, d_opt2, centerline_frenet,
347             centerline, headings, track_width, s_goal, s_next_1,
348             s_next_2,
349             U1_new, N, dt, d_g, d_u, d_t, d_next, d_c
350         )
351
352         # Check for convergence
353         delta_U1 = np.linalg.norm(U1_new - U1)
354         delta_U2 = np.linalg.norm(U2_new - U2)
355         if delta_U1 < epsilon and delta_U2 < epsilon:
356             break
357
358         U1, U2 = U1_new, U2_new
359
360         # Print cost value to check if its a saddle_point equilibrium
361         print("Cost for player 1: ", cost_J1)
362         print("Cost for player 2: ", cost_J2)
363
364     return U1[:, 0], U2[:, 0]
365
366 ## Solve the problem in Frenet Coordinates
367 def solve_zero_sum_qp_P1_frenet(

```

```

366     car1, car2,
367     d_opt1, d_opt2, centerline_frenet, kappa_interp, kappa_ref_1,
368     centerline, headings, track_width, s_goal, s_next_1, s_next_2,
369     P2_control_input,
370     N=10, dt=0.1, d_g=1.0, d_u=1.0, d_t=0.2, d_next=2.0, d_c=100.0
371 ):
372     # Init decision variables
373     n, m = 4, 2
374     X1 = cp.Variable((n, N+1))
375     U1 = cp.Variable((m, N))
376     U2 = P2_control_input
377     slack = cp.Variable(N)
378
379     # Initial state of the cars
380     s1_0, d1_0 = car1.get_frenet_coords(centerline, headings)
381     s2_0, d2_0 = car2.get_frenet_coords(centerline, headings)
382     state1_0 = car1.get_state()
383     state2_0 = car2.get_state()
384     dps1_0 = car1.compute_delta_psi(centerline, headings)
385     dps2_0 = car2.compute_delta_psi(centerline, headings)
386     x1_0 = np.array([s1_0, d1_0, state1_0[2], dps1_0])
387     x2_0 = np.array([s2_0, d2_0, state2_0[2], dps2_0])
388
389     constraints = [
390         X1[:, 0] == x1_0,
391     ]
392
393     cost = 0.0
394
395     # Assume horizon is short enough to linearize around x_0
396     A1, B1 = linearize_frenet_dynamics(x1_0, [0.0, 0.0], kappa_ref_1)
397     A1_d, B1_d = discretize_linear_dynamics(A1, B1, dt)
398
399     # Create interpolators locally
400     d_opt_interp1 = interp1d(centerline_frenet, d_opt1, kind='linear',
401                             fill_value='extrapolate')
402     d_opt_interp2 = interp1d(centerline_frenet, d_opt2, kind='linear',
403                             fill_value='extrapolate')
404
405     # Compute desired lateral offset
406     d_goal1 = d_opt_interp1(s1_0)
407
408     # Initialize state variables for car 2
409     x2_k = x2_0
410
411     for k in range(N):
412         # Constraints on dynamics
413         constraints += [X1[:, k+1] == A1_d @ X1[:, k] + B1_d @ U1[:, k]]
414
415         # Frenet coordinates for car 2

```



```

415     s2 = x2_k[0]
416     d2 = x2_k[1]
417
418     # Box constraints on speed, acceleration, steering and track
         width
419     constraints += [X1[3, k] <= car1.v_max, X1[3, k] >= 0.0]
420     constraints += [X1[1, k] <= track_width/2*0.95, X1[1, k] >= -
         track_width/2*0.95]
421     constraints += [U1[0, k] <= car1.a_max, U1[0, k] >= -car1.a_max]
422     constraints += [U1[1, k] <= car1.max_steering_angle, U1[1, k] >=
         -car1.max_steering_angle]
423
424     # Collision constraints
425     # rel = cp.vstack([X1[0,k] - s2, X1[1,k] - d2])
426     # Q = np.diag([1/2.0**2, 1/1.0**2])
427     # collision_expr = cp.quad_form(rel, Q)
428     # constraints += [collision_expr + slack[k] >= 1.0]
429     # constraints += [slack[k] >= 0]
430     # cost += d_c * slack[k]
431
432     # Update optimal trajectory goal
433     d_goal2 = d_opt_interp2(s2)
434
435     # Goal cost
436     cost += d_g * cp.square(s_goal - X1[0,k]) - d_g * cp.square(
         s_goal - s2)
437
438     # Optimal trajectory cost
439     cost += d_t * cp.square(d_goal1 - X1[1,k]) - d_t * cp.square(
         d_goal2 - d2)
440
441     # Input cost on acceleration
442     cost += d_u * cp.square(U1[0, k]) - d_u * cp.square(U2[0, k])
443
444     # Next waypoint cost
445     cost += d_next * cp.square(s_next_1 - X1[0,k]) - d_next * cp.
         square(s_next_2 - s2)
446
447     # Update state of car 2
448     x2_k = frenet_dynamics_discrete(x2_k, U2[:,k], kappa_interp, dt)
449
450     # Frenet coordinates for car 2 at last step
451     s2_N = x2_k[0]
452     d2_N = x2_k[1]
453
454     # Update optimal trajectory goal at last step
455     d_goal2 = d_opt_interp2(s2_N)
456
457     # Goal cost
458     cost += d_g * cp.square(s_goal - X1[0,N]) - d_g * cp.square(s_goal -
         s2_N)

```

```

459
460 # Optimal trajectory cost
461 cost += d_t * cp.square(d_goal1 - X1[1,N]) - d_t * cp.square(d_goal2
    - d2_N)
462
463 # Next waypoint cost
464 cost += d_next * cp.square(s_next_1 - X1[0,N]) - d_next * cp.square(
    s_next_2 - s2_N)
465
466 # Problem solving
467 prob = cp.Problem(cp.Minimize(cost), constraints)
468 prob.solve(solver=cp.OSQP)
469
470 # Check if solver succeeded
471 if prob.status not in [cp.OPTIMAL, cp.OPTIMAL_INACCURATE]:
472     print("[Warning] Car 1 QP solver failed:", prob.status)
473     return np.zeros((2, N))
474
475 return U1.value, prob.value
476
477 def solve_zero_sum_qp_P2_frenet(
478     car1, car2,
479     d_opt1, d_opt2, centerline_frenet, kappa_interp, kappa_ref_2,
480     centerline, headings, track_width, s_goal, s_next_1, s_next_2,
481     P1_control_input,
482     N=10, dt=0.1, d_g=1.0, d_u=1.0, d_t=0.2, d_next=2.0, d_c=100.0
483 ):
484     # Init decision variables
485     n, m = 4, 2
486     X2 = cp.Variable((n, N+1))
487     U2 = cp.Variable((m, N))
488     U1 = P1_control_input
489     slack = cp.Variable(N)
490
491     # Initial state of the cars
492     s1_0, d1_0 = car1.get_frenet_coords(centerline, headings)
493     s2_0, d2_0 = car2.get_frenet_coords(centerline, headings)
494     state1_0 = car1.get_state()
495     state2_0 = car2.get_state()
496     dps1_0 = car1.compute_delta_psi(centerline, headings)
497     dps2_0 = car2.compute_delta_psi(centerline, headings)
498     x1_0 = np.array([s1_0, d1_0, state1_0[2], dps1_0])
499     x2_0 = np.array([s2_0, d2_0, state2_0[2], dps2_0])
500
501     constraints = [
502         X2[:, 0] == x2_0,
503     ]
504
505     cost = 0.0
506
507     # Assume horizon is short enough to linearize around x_0

```

```

508 A2, B2 = linearize_frenet_dynamics(x2_0, [0.0, 0.0], kappa_ref_2)
509 A2_d, B2_d = discretize_linear_dynamics(A2, B2, dt)
510
511 # Create interpolators locally
512 d_opt_interp1 = interp1d(centerline_frenet, d_opt1, kind='linear',
513     fill_value='extrapolate')
514 d_opt_interp2 = interp1d(centerline_frenet, d_opt2, kind='linear',
515     fill_value='extrapolate')
516
517 # Compute desired lateral offset
518 d_goal2 = d_opt_interp1(s2_0)
519
520 # Initialize state variables for car 2
521 x1_k = x1_0
522
523 for k in range(N):
524
525     # Constraints on dynamics
526     constraints += [X2[:, k+1] == A2_d @ X2[:, k] + B2_d @ U2[:, k]]
527
528     # Frenet coordinates for car 1
529     s1 = x1_k[0]
530     d1 = x1_k[1]
531
532     # Box constraints on speed, acceleration, steering and track
533     # width
534     constraints += [X2[3, k] <= car2.v_max, X2[3, k] >= 0.0]
535     constraints += [X2[1, k] <= track_width/2*0.95, X2[1, k] >= -
536         track_width/2*0.95]
537     constraints += [U2[0, k] <= car2.a_max, U2[0, k] >= -car2.a_max]
538     constraints += [U2[1, k] <= car2.max_steering_angle, U2[1, k] >=
539         -car2.max_steering_angle]
540
541     # Collision constraints
542     # rel = cp.vstack([X2[0,k] - s1, X2[1,k] - d1])
543     # Q = np.diag([1/2.0**2, 1/1.0**2])
544     # collision_expr = cp.quad_form(rel, Q)
545     # constraints += [collision_expr + slack[k] >= 1.0]
546     # constraints += [slack[k] >= 0]
547     # cost += d_c * slack[k]
548
549     # Update optimal trajectory goal
550     d_goal1 = d_opt_interp2(s1)
551
552     # Goal cost
553     cost += d_g * cp.square(s_goal - s1) - d_g * cp.square(s_goal -
554         X2[0,k])
555
556     # Optimal trajectory cost
557     cost += d_t * cp.square(d_goal1 - d1) - d_t * cp.square(d_goal2
558         - X2[1,k])
    
```

```

552     # Input cost on acceleration
553     cost += d_u * cp.square(U1[0, k]) - d_u * cp.square(U2[0, k])
554
555     # Next waypoint cost
556     cost += d_next * cp.square(s_next_1 - s1) - d_next * cp.square(
557         s_next_2 - X2[0, k])
558
559     # Update state of car 2
560     x1_k = frenet_dynamics_discrete(x1_k, U1[:, k], kappa_interp, dt)
561
562     # Frenet coordinates for car 2 at last step
563     s1_N = x1_k[0]
564     d1_N = x1_k[1]
565
566     # Update optimal trajectory goal at last step
567     d_goal1 = d_opt_interp2(s1_N)
568
569     # Goal cost
570     cost += d_g * cp.square(s_goal - s1_N) - d_g * cp.square(s_goal - X2
571         [0, N])
572
573     # Optimal trajectory cost
574     cost += d_t * cp.square(d_goal1 - d1_N) - d_t * cp.square(d_goal2 -
575         X2[1, N])
576
577     # Next waypoint cost
578     cost += d_next * cp.square(s_next_1 - s1_N) - d_next * cp.square(
579         s_next_2 - X2[0, N])
580
581     # Problem solving
582     prob = cp.Problem(cp.Minimize(-cost), constraints)
583     prob.solve(solver=cp.OSQP)
584
585     # Check if solver succeeded
586     if prob.status not in [cp.OPTIMAL, cp.OPTIMAL_INACCURATE]:
587         print("[Warning] Car 2 QP solver failed:", prob.status)
588         return np.zeros((2, N))
589
590     return U2.value, prob.value
591
592 def zero_sum_best_response_frenet(
593     car1, car2,
594     d_opt1, d_opt2, centerline_frenet, kappa_interp, kappa_ref_1,
595     kappa_ref_2,
596     centerline, headings, track_width, s_goal, s_next_1, s_next_2,
597     N=10, dt=0.1, d_g=10.0, d_u=1.0, d_t=30.0, d_next=20.0, d_c=100.0,
598     max_iters=5, epsilon=1e-6
599 ):
600     np.random.seed(42)

```

```

598     m = 2  # dimension of control input
599     U1 = np.zeros((m, N))
600     U2 = np.random.randn(m, N)
601     U2[0, :] *= car2.a_max
602     U2[1, :] *= car2.max_steering_angle
603
604     for i in range(max_iters):
605         print(f"  Iter {i+1}/{max_iters}")
606
607         # Minimize cost for Car 1 given U2
608         U1_new, cost_J1 = solve_zero_sum_qp_P1_frenet(
609             car1, car2,
610             d_opt1, d_opt2, centerline_frenet, kappa_interp, kappa_ref_1
611             ,
612             centerline, headings, track_width, s_goal, s_next_1,
613             s_next_2,
614             U2, N, dt, d_g, d_u, d_t, d_next, d_c
615         )
616
617         # Step 2: Maximize cost for Car 2 given U1
618         U2_new, cost_J2 = solve_zero_sum_qp_P2_frenet(
619             car1, car2,
620             d_opt1, d_opt2, centerline_frenet, kappa_interp, kappa_ref_2
621             ,
622             centerline, headings, track_width, s_goal, s_next_1,
623             s_next_2,
624             U1_new, N, dt, d_g, d_u, d_t, d_next, d_c
625         )
626
627         # Check if solution converged
628         delta_U1 = np.linalg.norm(U1_new - U1)
629         delta_U2 = np.linalg.norm(U2_new - U2)
630         if delta_U1 < epsilon and delta_U2 < epsilon:
631             break
632
633         U1, U2 = U1_new, U2_new
634
635         # Print cost value to check if its a saddle_point equilibrium
636         print("Cost for player 1: ", cost_J1)
637         print("Cost for player 2: ", -cost_J2)
638
639     return U1[:, 0], U2[:, 0], cost_J1, -1*cost_J2

```

## Appendix I. Simulation Code

simulation.py

```

1  import math
2  import numpy as np
3  import matplotlib.pyplot as plt

```

```

4 from scipy.interpolate import interp1d
5 from racetrack import generate_track
6 from optimal_trajectory import generate_optimal_traj
7 from car import Car
8 from ZS_Controller import zero_sum_best_response_frenet
9 from frenet import compute_curvature
10
11 def resample_centerline(centerline, num_points=500):
12     '''Function used to sample additional points along the centerline
13     since trackgen only returns the minimal amount of points
14     necessary to describe the track'''
15     ds = np.hypot(np.diff(centerline[:, 0]), np.diff(centerline[:, 1]))
16     s = np.insert(np.cumsum(ds), 0, 0.0)
17     s_uniform = np.linspace(0, s[-1], num_points)
18     interp_x = interp1d(s, centerline[:, 0], kind='linear')
19     interp_y = interp1d(s, centerline[:, 1], kind='linear')
20     centerline_dense = np.stack([interp_x(s_uniform), interp_y(s_uniform)
21     ], axis=1)
22     return centerline_dense, s_uniform
23
24 # Generate the track
25 num_points = 1000
26 Length, Width = 100, 6
27 track = generate_track(Length, Width)
28 track.plot()
29 centerline_raw = np.column_stack((track.xm, track.ym))
30
31 # Remove duplicate (x,y) points from the centerline (trackgen can
32 duplicate points sometimes and they can cause numerical errors)
33 _, unique_indices = np.unique(centerline_raw, axis=0, return_index=True)
34 centerline_raw = centerline_raw[np.sort(unique_indices)]
35
36 # Resample the centerline
37 centerline, centerline_frenet = resample_centerline(centerline_raw,
38 num_points)
39
40 # Compute centerline headings
41 diffs = np.diff(centerline, axis=0)
42 headings = np.arctan2(diffs[:, 1], diffs[:, 0])
43 headings = np.append(headings, headings[-1])
44
45 # Initialize cars
46 g = 9.81
47 car1 = Car([track.xm[0], track.ym[0]+2, headings[0], 5.0], v_max=65.0,
48 a_max=10.0, a_lat_max=1*g)
49 car2 = Car([track.xm[0], track.ym[0]-2, headings[0], 5.0], v_max=65.0,
50 a_max=10.0, a_lat_max=2*g)
51
52 # Generate the optimal trajectory for each car given the track
53 d_opt1 = generate_optimal_traj(centerline_frenet, track.width, car1.
54 a_lat_max)

```

## ZERO-SUM RACING GAME

```

47 d_opt2 = generate_optimal_traj(centerline_frenet, track.width, car2.
    a_lat_max)
48
49 # Compute the curvature along the centerline (Artificial, only works for
    the U shaped track)
50 kappa = np.zeros_like(headings)
51 straight_angles = [0, np.pi, -np.pi]
52 tolerance = 1e-2
53 # Curvature is zero for straight lines (heading is 0 or pi or -pi)
54 curved_mask = ~np.isclose(headings % (2*np.pi), straight_angles[0], atol
    =tolerance) & \
55     ~np.isclose(headings % (2*np.pi), straight_angles[1], atol
    =tolerance) & \
56     ~np.isclose(headings % (2*np.pi), straight_angles[2], atol
    =tolerance)
57
58 # kappa[curved_mask] = 1.0 / 40.0 # curvature of a semi-circle is
    simply 1/R
59
60 # Curvature interpolation function used in the discrete dynamics
61 kappa_interp = interp1d(centerline_frenet, kappa, kind='linear',
    bounds_error=False, fill_value="extrapolate")
62
63 # Useful prints for debugging
64 # print("centerline :", centerline)
65 # print("headings :", headings)
66 # print("centerline_frenet :", centerline_frenet)
67 # print("dopt1 :", d_opt1)
68 # print("dopt2 :", d_opt2)
69 # print("curvature:", kappa)
70
71 ## Simulation
72 # Simulation parameters init
73 s_goal = centerline_frenet[-1]
74 N = 20
75 dt = 0.05
76 T = 40
77
78 # Cost lists init
79 J1 = []
80 J2 = []
81
82 # Loop
83 for t in range(T):
84     print(f"Step {t+1}/{T}")
85
86     # Get current Frenet coordinates
87     s1, _ = car1.get_frenet_coords(centerline, headings)
88     s2, _ = car2.get_frenet_coords(centerline, headings)
89
90     # Find the closest index in the centerline

```

```

91     idx_closest_1 = np.argmin(np.abs(centerline_frenet - s1))
92     idx_closest_2 = np.argmin(np.abs(centerline_frenet - s2))
93
94     # Move couple steps forward
95     idx_next_1 = min(idx_closest_1 + int(math.floor(num_points/10)), len
96                     (centerline_frenet) - 1)
97     idx_next_2 = min(idx_closest_2 + int(math.floor(num_points/10)), len
98                     (centerline_frenet) - 1)
99
100    # Next reference point in s
101    s_next_1 = centerline_frenet[idx_next_1]
102    s_next_2 = centerline_frenet[idx_next_2]
103
104    # Compute centerline curvature at current car position
105    kappa_ref_1 = kappa_interp(s1)
106    kappa_ref_2 = kappa_interp(s2)
107
108    # Compute optimal input for both cars using the iterative best
109    # response approach
110    u1, u2, cost_J1, cost_J2 = zero_sum_best_response_frenet(
111        car1, car2,
112        d_opt1, d_opt2, centerline_frenet, kappa_interp, kappa_ref_1,
113        kappa_ref_2,
114        centerline, headings, Width, s_goal, s_next_1, s_next_2,
115        N=N, dt=dt
116    )
117    print("u1: ", u1)
118    print("u2: ", u2)
119    car1.set_control(*u1)
120    car2.set_control(*u2)
121    car1.update()
122    car2.update()
123    J1.append(cost_J1)
124    J2.append(cost_J2)
125
126    # Statistics on cost to assess if saddle point equilibrium is reached or
127    # not
128    J1_array = np.array(J1)
129    J2_array = np.array(J2)
130
131    # Compute absolute difference
132    diff = np.abs(J1_array - J2_array)
133
134    # Compute statistics
135    mean_diff = np.mean(diff)
136    median_diff = np.median(diff)
137    std_diff = np.std(diff)
138    min_diff = np.min(diff)
139    max_diff = np.max(diff)
140
141    # Print statistics

```



```

137 print("Statistics of |J1 - J2|:")
138 print(f"Mean      : {mean_diff:.4f}")
139 print(f"Median    : {median_diff:.4f}")
140 print(f"Std Dev:   {std_diff:.4f}")
141 print(f"Min       : {min_diff:.4f}")
142 print(f"Max       : {max_diff:.4f}")
143
144 # Create box plot
145 plt.figure(figsize=(6, 5))
146 plt.boxplot(diff, vert=True, patch_artist=True)
147 plt.title("Distribution of |J1 - J2|")
148 plt.ylabel("Absolute Cost Difference")
149 plt.grid(True)
150 plt.show()
151
152 # Plotting results
153 h1 = np.array(car1.history)
154 h2 = np.array(car2.history)
155
156 plt.figure(figsize=(10, 8))
157 plt.plot(track.xm, track.ym, 'k--', label='Centerline')
158 plt.plot(track.xb1, track.yb1, 'g--', label='Inner Boundary')
159 plt.plot(track.xb2, track.yb2, 'r--', label='Outer Boundary')
160 plt.plot(h1[:, 0], h1[:, 1], 'b-', label='Car 1 Trajectory')
161 plt.plot(h2[:, 0], h2[:, 1], 'orange', label='Car 2 Trajectory')
162 plt.ylim(-10, 10)
163 plt.legend()
164 plt.title("Zero-Sum Racing Simulation")
165 plt.xlabel("X [m]")
166 plt.ylabel("Y [m]")
167 plt.grid(True)
168 plt.show()

```